

VISTORA ASSIGNMENT

PRESENTED BY ASHUTOSH SAHOO



Introduction to Feature Engineering

Feature engineering is the activity of developing, choosing, or manipulating input variables (features) to enhance the performance of a machine learning model. It increases accuracy, decreases complexity, and enables models to learn patterns more effectively, thus being important in constructing efficient and effective ML systems.

Types of Feature Engineering

Binning – Groups continuous data into discrete intervals (e.g., age groups) to reduce noise and overfitting.

Feature Interaction – Combines features (e.g., multiplying or summing) to reveal relationships that improve predictive power.

Normalization/Scaling – Adjusts numerical features to a common scale (e.g., Min-Max, Z-score) to prevent bias from varying magnitudes.

Encoding – Converts categorical data into numerical form (e.g., One-Hot, Label Encoding) for model compatibility.

Time-Based Aggregations – Extracts trends from timestamps (e.g., daily averages, rolling stats) to capture temporal patterns



Using Snowflake for Data Storage & Processing

How Snowflake is used for storing structured and semi-structured data.

Snowflake keeps structured data (tables, CSVs) in standard relational formats and semi-structured data (JSON, XML, Parquet) in VARIANT columns, which automatically parse nested fields. Schema-on-read enables querying semi-structured data through SQL without pre-transformations. Inbuilt functions such as FLATTEN extract nested elements, which facilitates effortless analytics over mixed data types without compromising performance and scalability.



Example of SQL queries that extract and preprocess data in Snowflake

-- 1. Extract & clean structured data (e.g., sales table)

```
SELECT
  ORDER_ID,
  CUSTOMER_ID,
  TO_DATE(ORDER_DATE, 'YYYY-MM-DD') AS
  FORMATTED_DATE, -- Date normalization
  NULLIF(TOTAL_AMOUNT, 0) AS ADJUSTED_AMOUNT --
  Handle zeros
FROM SALES
WHERE STATUS = 'COMPLETED';
```

-- 2. Parse semi-structured JSON logs (e.g., event data)

```
SELECT
  LOG_ID,
  EVENT_DATA:user_id::STRING AS USER_ID, -- Extract
  JSON field
  EVENT_DATA:timestamp::TIMESTAMP AS EVENT_TIME,
  LOWER(EVENT_DATA:action::STRING) AS ACTION --
  Normalize text
FROM EVENT_LOGS
WHERE EVENT_TIME > '2024-01-01';
```


How Snowflake integrates with ML pipelines.

Snowflake naturally fits into ML pipelines by facilitating direct access to data through Snowpark (for Python/Scala) or through connectors such as Spark. In-database model training is supported through Snowflake ML, and external tools (e.g., TensorFlow, SageMaker) can query processed data through SQL. Efficient feature engineering is ensured through Snowflake's scalability, and its secure data sharing makes collaboration between teams easier, all the way from data preparation to deployment.



Feature Store Concepts

A Feature Store is a centralized system used to store, manage, and serve machine learning features reliably throughout training and production. It promotes feature consistency, prevents repeated preprocessing, and provides real-time access for model inference. Maintaining a single source of truth for features makes it speed up ML development, lower errors, and enhance collaboration among data scientists and engineers, enabling model deployment to be faster and more dependable

Compare different feature stores

AWS SageMaker Feature Store excels in real-time and batch feature serving, tightly integrated with AWS ML services. Snowflake Feature Store leverages Snowflake's SQL engine for seamless feature management within its data cloud, ideal for structured/semi-structured data. Databricks Feature Store integrates with MLflow and Delta Lake, optimizing feature sharing across Spark workflows. While AWS focuses on end-to-end ML, Snowflake emphasizes SQL-native features, and Databricks unifies feature engineering with big data processing, catering to distinct ecosystem needs.

Implementation And Practical Task

Connect to Snowflake, extract some sample data,
and perform feature engineering.

```
USE DATABASE SNOWFLAKE_SAMPLE_DATA;  
SHOW SCHEMAS;  
USE DATABASE SNOWFLAKE_SAMPLE_DATA;  
USE SCHEMA TPCH_SF1;  
SHOW TABLES;
```

```
SELECT * FROM CUSTOMER LIMIT 10;  
SELECT * FROM ORDERS LIMIT 1000;  
SELECT * FROM LINEITEM LIMIT 10;
```

	# O_ORDERKEY	# O_CUSTKEY	A O_ORDERSTATUS	# O_TOTALPRICE	⌚ O_ORDERDATE	A O_ORDERPRIORITY	A O_CLERK
1	3000001	145618	F	30175.88	1992-12-17	4-NOT SPECIFIED	Clerk#0000
2	3000002	1481	O	297999.63	1995-07-28	1-URGENT	Clerk#0000
3	3000003	127432	O	345438.38	1997-11-04	5-LOW	Clerk#0000
4	3000004	47423	O	135965.53	1996-06-13	4-NOT SPECIFIED	Clerk#0000
5	3000005	84973	F	209937.09	1992-09-12	5-LOW	Clerk#0000
6	3000006	135136	O	140186.32	1996-09-26	1-URGENT	Clerk#0000
7	3000007	78841	F	298655.07	1992-04-13	5-LOW	Clerk#0000
8	3000032	124576	F	175973.90	1992-03-02	1-URGENT	Clerk#0000
9	3000033	30247	F	4635.38	1993-11-10	1-URGENT	Clerk#0000
10	3000034	5498	F	348308.79	1992-04-21	1-URGENT	Clerk#0000

The Normalizing step of Feature Engineering

-- Normalize O_TOTALPRICE Column

```
WITH stats AS (  
  SELECT  
    MIN(O_TOTALPRICE) AS min_price,  
    MAX(O_TOTALPRICE) AS max_price  
  FROM ORDERS  
)  
SELECT  
  O_ORDERKEY,  
  O_TOTALPRICE,  
  (O_TOTALPRICE - stats.min_price) / (stats.max_price - stats.min_price) AS totalprice_normalized  
FROM ORDERS, stats  
LIMIT 10;
```

	# O_ORDERKEY	# O_TOTALPRICE	# TOTALPRICE_NORMA
1	5400001	270576.60	0.48648
2	5400002	216696.22	0.38929
3	5400003	191044.99	0.34303
4	5400004	263505.65	0.47372
5	5400005	117459.27	0.21030
6	5400006	84588.56	0.15102
7	5400007	50890.64	0.09024
8	5400032	24285.54	0.04225
9	5400033	181139.61	0.32516
10	5400034	54996.61	0.09764

Store the features in a Feature Store

```
CREATE DATABASE IF NOT EXISTS FEATURE_DB;  
USE DATABASE FEATURE_DB;  
CREATE SCHEMA IF NOT EXISTS FEATURE_SCHEMA;  
USE SCHEMA FEATURE_SCHEMA;
```

```
CREATE OR REPLACE TABLE ORDER_FEATURES AS  
WITH stats AS (  
  SELECT MIN(O_TOTALPRICE) AS min_price, MAX(O_TOTALPRICE) AS max_price  
  FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF1.ORDERS  
)  
cust_first_order AS (  
  SELECT O_CUSTKEY, MIN(O_ORDERDATE) AS first_order_date  
  FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF1.ORDERS  
  GROUP BY O_CUSTKEY  
)  
SELECT  
  O.O_ORDERKEY,  
  O.O_CUSTKEY,  
  O.O_ORDERDATE,
```

-- Date Features

```
EXTRACT(YEAR FROM O.O_ORDERDATE) AS order_year,  
EXTRACT(MONTH FROM O.O_ORDERDATE) AS order_month,  
EXTRACT(DAY FROM O.O_ORDERDATE) AS order_day,  
DAYOFWEEK(O.O_ORDERDATE) AS order_dayofweek,  
WEEKOFYEAR(O.O_ORDERDATE) AS order_weekofyear,
```

-- Encoded ORDERSTATUS

```
CASE WHEN O.O_ORDERSTATUS = 'F' THEN 1 ELSE 0 END AS is_filled,  
CASE WHEN O.O_ORDERSTATUS = 'O' THEN 1 ELSE 0 END AS is_open,  
CASE WHEN O.O_ORDERSTATUS = 'P' THEN 1 ELSE 0 END AS is_pending,
```

-- Encoded ORDERPRIORITY

```
CASE WHEN O.O_ORDERPRIORITY = '1-URGENT' THEN 1 ELSE 0 END AS is_urgent,  
CASE WHEN O.O_ORDERPRIORITY = '2-HIGH' THEN 1 ELSE 0 END AS is_high,  
CASE WHEN O.O_ORDERPRIORITY = '3-MEDIUM' THEN 1 ELSE 0 END AS is_medium,  
CASE WHEN O.O_ORDERPRIORITY = '4-NOT SPECIFIED' THEN 1 ELSE 0 END AS is_not_specified,
```

-- Normalized Price

```
O.O_TOTALPRICE,  
(O.O_TOTALPRICE - stats.min_price) / NULLIF((stats.max_price - stats.min_price), 0) AS totalprice_normalized,
```

-- Days since first order

```
DATEDIFF('day', cust_first_order.first_order_date, O.O_ORDERDATE) AS days_since_first_order
```

```
FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF1.ORDERS O  
JOIN cust_first_order ON O.O_CUSTKEY = cust_first_order.O_CUSTKEY,  
stats;
```


SELECT * FROM FEATURE_DB.FEATURE_SCHEMA.ORDER_FEATURES LIMIT 10;

DESC TABLE FEATURE_DB.FEATURE_SCHEMA.ORDER_FEATURES;

ResultsChart

	# O_ORDERKEY	# O_CUSTKEY	🕒 O_ORDERDATE	# ORDER_YEAR	# ORDER_MONTH	# ORDER_DAY	# ORDER_DAYOFWEEK
1	2400001	21445	1995-04-03	1995	4	3	
2	2400002	104548	1994-08-25	1994	8	25	
3	2400003	8512	1996-07-09	1996	7	9	
4	2400004	85286	1996-09-30	1996	9	30	
5	2400005	147364	1994-05-03	1994	5	3	
6	2400006	95339	1995-11-14	1995	11	14	
7	2400007	46249	1993-07-10	1993	7	10	
8	2400032	130357	1994-08-16	1994	8	16	
9	2400033	138706	1995-02-01	1995	2	1	
10	2400034	79036	1994-11-09	1994	11	9	

Query Details

Query duration85ms

Rows10

Query ID01bc6f35-3201-a0e1-0...

Show more

O_ORDERKEY

O_CUSTKEY

Retrieve features and link to the Python Notebook to use Features

```
import snowflake.connector
import pandas as pd

# Establish connection
conn = snowflake.connector.connect(
    user='AshutoshSahoo1234',
    password='AshutoshSahoo2025',
    account='IZDSRLY-GR41231',
    warehouse='COMPUTE_WH',
    database='FEATURE_DB',
    schema='FEATURE_SCHEMA'
)

try:
    # Query the feature table
    query = "SELECT * FROM ORDER_FEATURES"
    cursor = conn.cursor()
    cursor.execute(query)

    # Load into a DataFrame
    df = pd.DataFrame(cursor.fetchall(), columns=[col[0] for col in cursor.description])

finally:
    # Always close the cursor and connection
    cursor.close()
    conn.close()

# Now df contains your features and can be used for training
print(df.head())
```

	O_ORDERKEY	O_CUSTKEY	O_ORDERDATE	ORDER_YEAR	ORDER_MONTH	ORDER_DAY	\
0	1	36901	1996-01-02	1996	1	2	
1	2	78002	1996-12-01	1996	12	1	
2	3	123314	1993-10-14	1993	10	14	
3	4	136777	1995-10-11	1995	10	11	
4	5	44485	1994-07-30	1994	7	30	

	ORDER_DAYOFWEEK	ORDER_WEEKOFYEAR	IS_FILLED	IS_OPEN	IS_PENDING	\
0	2	1	0	1	0	
1	0	48	0	1	0	
2	4	41	1	0	0	
3	3	41	0	1	0	
4	6	30	1	0	0	

	IS_URGENT	IS_HIGH	IS_MEDIUM	IS_NOT_SPECIFIED	O_TOTALPRICE	\
0	0	0	0	0	173665.47	
1	1	0	0	0	46929.18	
2	0	0	0	0	193846.25	
3	0	0	0	0	32151.78	
4	0	0	0	0	144659.20	

	TOTALPRICE_NORMALIZED	DAYS_SINCE_FIRST_ORDER
0	0.31168688	1380

Thank you